

```
from google.colab import drive
```

```
drive.mount('/content/drive')
```

```
↳ Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True)
```

```
%cd 'drive'
```

```
↳ /content/drive
```

```
%cd 'My Drive/tehtable'
```

```
↳ /content/drive/My Drive/tehtable
```

```
from cifar100 import *
```

```
from torchvision import transforms
stats = ((0.4914, 0.4822, 0.4465), (0.2023, 0.1994, 0.2010))
transform_test=transforms.Compose([
    transforms.RandomCrop(32, padding=4, padding_mode='reflect'),
    transforms.RandomHorizontalFlip(),
    transforms.ToTensor(),
    transforms.Normalize(*stats,inplace=True)
])
```

```
transform_test1 = transforms.Compose([transforms.ToTensor(), transforms.Normalize(*stats)])
```

```
train=CIFAR100(root='data',download=True,transform=transform_test)
test=CIFAR100(root='data',download=True,train=False,transform=transform_test1)
```

```

↳ Files already downloaded and verified
Files already downloaded and verified

```

```
print(train)
```

```

↳ Dataset CIFAR100
  Number of datapoints: 50000
  Split: train
  Root Location: data
  Transforms (if any): Compose(
    RandomCrop(size=(32, 32), padding=4)
    RandomHorizontalFlip(p=0.5)
    ToTensor()
    Normalize(mean=(0.4914, 0.4822, 0.4465), std=(0.2023, 0.1994, 0.201))
  )
  Target Transforms (if any): None

```

```
img,label,index=train.__getitem__(65)
```

```
label
```

```
↳ 35
```

```
test.__len__()
```

```
↳ 10000
```

```
train.__repr__()
```

```

↳ 'Dataset CIFAR100\n  Number of datapoints: 50000\n  Split: train\n  Root Location: data\n  Transforms (if any): Compos
e(\n    RandomCrop(size=(32, 32), padding=4)\n    RandomHorizontalFlip(p=0.
5)\n    ToTensor()\n    Normalize(mean=(0.4914, 0.4822, 0.4465), std=(0.202
3, 0.1994, 0.201))\n  )\n  Target Transforms (if any): None'

```

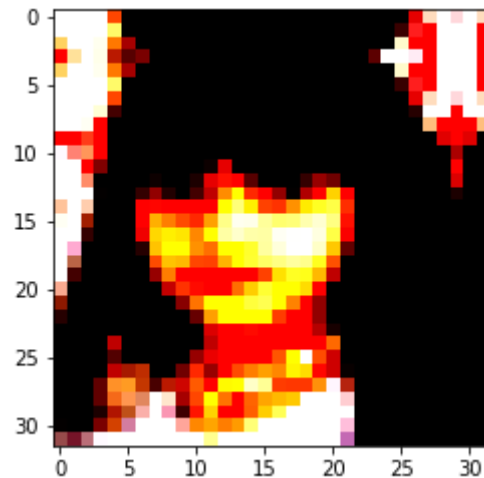
```

import matplotlib.pyplot as plt
%matplotlib inline

```

```
plt.imshow(img.permute(1,2,0))
print(label)
```

☞ Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).



```
import torch
from torch.utils.data import random_split
```

```
random_seed=42
torch.manual_seed(random_seed)
```

☞ <torch._C.Generator at 0x7fed7f61ab90>

```
#Splitting into Train and Validation Sets
val_size=5000
train_size=len(train)-val_size
```

```
train_ds,val_ds=random_split(train,[train_size,val_size])
```

```
len(train_ds),len(val_ds)
```

↗ (45000, 5000)

```
#Loading the images in batches
from torch.utils.data.dataloader import DataLoader
batch_size=128

train_dl=DataLoader(train_ds,batch_size,shuffle=True,num_workers=4,pin_memory=True)
val_dl=DataLoader(val_ds,batch_size*2,num_workers=4,pin_memory=True)

from torchvision.utils import make_grid

def show_batch(dl):
    for images, labels,_ in dl:
        fig, ax = plt.subplots(figsize=(12, 6))
        ax.set_xticks([]); ax.set_yticks([])
        ax.imshow(make_grid(images, nrow=16).permute(1, 2, 0))
        break

show_batch(train_dl)
```

↗

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).



#Building the model



```
import torch.nn as nn
import torch.nn.functional as F
import pandas as pd
```



```
def accuracy(outputs, labels):
    _, preds = torch.max(outputs, dim=1)
    return torch.tensor(torch.sum(preds == labels).item() / len(preds))

class ImageClassificationBase(nn.Module):
    def training_step(self, batch):
        images, labels, _ = batch
        out = self(images)           # Generate predictions
        loss = F.cross_entropy(out, labels) # Calculate loss
        return loss

    def validation_step(self, batch):
        images, labels, _ = batch
        out = self(images)           # Generate predictions
        loss = F.cross_entropy(out, labels) # Calculate loss
        acc = accuracy(out, labels)     # Calculate accuracy
        return {'val_loss': loss.detach(), 'val_acc': acc}

    def validation_epoch_end(self, outputs):
        batch_losses = [x['val_loss'] for x in outputs]
        epoch_loss = torch.stack(batch_losses).mean() # Combine losses
        batch_accs = [x['val_acc'] for x in outputs]
        epoch_acc = torch.stack(batch_accs).mean()    # Combine accuracies
        return {'val_loss': epoch_loss.item(), 'val_acc': epoch_acc.item()}
```

```

def epoch_end(self, epoch, result):
    print("Epoch [{}], last_lr: {:.5f}, train_loss: {:.4f}, val_loss: {:.4f}, val_acc: {:.4f}".format(
        epoch, result['lrs'][-1], result['train_loss'], result['val_loss'], result['val_acc']))

def conv_block(in_channels, out_channels, pool=False):
    layers = [nn.Conv2d(in_channels, out_channels, kernel_size=3, padding=1),
              nn.BatchNorm2d(out_channels),
              nn.ReLU(inplace=True)]
    if pool: layers.append(nn.MaxPool2d(2))
    return nn.Sequential(*layers)

class ResNet9(ImageClassificationBase):
    def __init__(self, in_channels, num_classes):
        super().__init__()

        self.conv1 = conv_block(in_channels, 64)
        self.conv2 = conv_block(64, 128, pool=True)
        self.res1 = nn.Sequential(conv_block(128, 128), conv_block(128, 128))

        self.conv3 = conv_block(128, 256, pool=True)
        self.conv4 = conv_block(256, 512, pool=True)
        self.res2 = nn.Sequential(conv_block(512, 512), conv_block(512, 512))

        self.classifier = nn.Sequential(nn.MaxPool2d(4),
                                         nn.Flatten(),
                                         nn.Linear(512, num_classes))

    def forward(self, xb):
        out = self.conv1(xb)
        out = self.conv2(out)
        out = self.res1(out) + out
        out = self.conv3(out)
        out = self.conv4(out)
        out = self.res2(out) + out
        out = self.classifier(out)
        return out

```

```
model = ResNet9(3, 100)  
model
```



```

ResNet9(
  (conv1): Sequential(
    (0): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (2): ReLU(inplace=True)
  )
  (conv2): Sequential(
    (0): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (2): ReLU(inplace=True)
    (3): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  )
  (res1): Sequential(
    (0): Sequential(
      (0): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (2): ReLU(inplace=True)
    )
  )
)

def get_default_device():

    if torch.cuda.is_available():
        return torch.device('cuda')
    else:
        return torch.device('cpu')

def to_device(data, device):

    if isinstance(data, (list,tuple)):
        return [to_device(x, device) for x in data]
    return data.to(device, non_blocking=True)

class DeviceDataLoader():

    def __init__(self, dl, device):
        self.dl = dl
        self.device = device

    def __iter__(self):

```



```

        for b in self.dl:
            yield to_device(b, self.device)

    def __len__(self):
        return len(self.dl)

    (0): MaxPool2d(kernel_size=4, stride=4, padding=0, dilation=1, ceil_mode=False)
device = get_default_device()
device

↳ device(type='cuda')

train_dl = DeviceDataLoader(train_dl, device)
val_dl = DeviceDataLoader(val_dl, device)
to_device(model, device);

@torch.no_grad()
def evaluate(model, val_loader):
    model.eval()
    outputs = [model.validation_step(batch) for batch in val_loader]
    return model.validation_epoch_end(outputs)

def get_lr(optimizer):
    for param_group in optimizer.param_groups:
        return param_group['lr']

def fit_one_cycle(epochs, max_lr, model, train_loader, val_loader,
                  weight_decay=0, grad_clip=None, opt_func=torch.optim.SGD):
    torch.cuda.empty_cache()
    history = []

    # Set up custom optimizer with weight decay
    optimizer = opt_func(model.parameters(), max_lr, weight_decay=weight_decay)
    # Set up one-cycle learning rate scheduler
    sched = torch.optim.lr_scheduler.OneCycleLR(optimizer, max_lr, epochs=epochs,
                                                  steps_per_epoch=len(train_loader))

```

```

for epoch in range(epochs):
    # Training Phase
    model.train()
    train_losses = []
    lrs = []
    for batch in train_loader:
        loss = model.training_step(batch)
        train_losses.append(loss)
        loss.backward()

        # Gradient clipping
        if grad_clip:
            nn.utils.clip_grad_value_(model.parameters(), grad_clip)

        optimizer.step()
        optimizer.zero_grad()

        # Record & update learning rate
        lrs.append(get_lr(optimizer))
        sched.step()

    # Validation phase
    result = evaluate(model, val_loader)
    result['train_loss'] = torch.stack(train_losses).mean().item()
    result['lrs'] = lrs
    model.epoch_end(epoch, result)
    history.append(result)
return history

```

```
model = to_device(ResNet9(3,100), device)
```

```
history=[evaluate(model, val_dl)]
history
```



```

[{'loss': 0.011014062212725485, 'val_loss': 4.6066484451202045}]
epochs = 5
max_lr = 0.01
grad_clip = 0.1
weight_decay = 1e-4
opt_func = torch.optim.Adam

history += fit_one_cycle(epochs, max_lr, model, train_dl, val_dl,
                        grad_clip=grad_clip,
                        weight_decay=weight_decay,
                        opt_func=opt_func)


↳ Epoch [0], last_lr: 0.00759, train_loss: 1.4882, val_loss: 2.5325, val_acc: 0.3607
Epoch [1], last_lr: 0.00950, train_loss: 1.9473, val_loss: 2.6184, val_acc: 0.3727
Epoch [2], last_lr: 0.00611, train_loss: 1.7668, val_loss: 2.0248, val_acc: 0.4639
Epoch [3], last_lr: 0.00188, train_loss: 1.4133, val_loss: 1.4205, val_acc: 0.6071
Epoch [4], last_lr: 0.00000, train_loss: 1.0351, val_loss: 1.2560, val_acc: 0.6434

def plot_accuracies(history):
    accuracies = [x['val_acc'] for x in history]
    plt.plot(accuracies, '-x')
    plt.xlabel('epoch')
    plt.ylabel('accuracy')
    plt.title('Accuracy vs. No. of epochs');

plot_accuracies(history)


```

↳



```
def plot_losses(history):
    train_losses = [x.get('train_loss') for x in history]
    val_losses = [x['val_loss'] for x in history]
    plt.plot(train_losses, '-bx')
    plt.plot(val_losses, '-rx')
    plt.xlabel('epoch')
    plt.ylabel('loss')
    plt.legend(['Training', 'Validation'])
    plt.title('Loss vs. No. of epochs');
```

plot_losses(history)



```
test_dl=DataLoader(test,batch_size^2,num_workers=4,pin_memory=True)  
test_dl = DeviceDataLoader(test_dl, device)
```

```
evaluate(model,test_dl)
```

```
↳ {'val_acc': 0.6639648675918579, 'val_loss': 1.1835086345672607}
```

```
train_on_gpu=torch.cuda.is_available()
```

```
predicted_values=[]  
model.eval()  
for data,target,_ in test_dl:  
    output=model(data)  
    _,pred=torch.max(output,1)  
    p=pred.cpu().numpy()  
    t=p.tolist()  
    predicted_values.append(t)
```

```
print(len(predicted_values))
```

```
↳ 40
```

```
flat_list = []  
for sublist in predicted_values:  
    for item in sublist:  
        flat_list.append(item)
```

```
arr=np.arange(0,10000)
```

```
arr=pd.Series(arr)
```

```
arr1=arr.to_frame()

arr2=pd.DataFrame(flat_list)

df=pd.concat([arr1,arr2],ignore_index=True,axis=1)

df.to_csv("submission.csv",header=["Id","Category"],index=False)
```